

The system already exists.

A batch rendering engine built, tested and running — prompt file in, numbered images out, one zipped folder per batch. Not a plan to build one.

SDXL	Lightning 4-step
700	Per batch, in order
15+	Batches per day
~2s	Per image on A10G

What you get, exactly as specified

You send a TXT file with ~700 prompts, one per line. You get back a single folder of 1.jpg through 700.jpg, in your exact prompt order, zipped.

Prompt order is guaranteed structurally, not by hope. Each image is named from its line number in your file, so prompt 437 always becomes 437.jpg. Blank lines and Windows line endings can't shift the numbering. Before a batch is zipped, every expected index is verified present and decodable — if anything is missing, the run tells me rather than quietly shipping 699 images.

The stack

Layer	Choice	Why
Model	SDXL-Lightning 4-step on SDXL base 1.0	Near-SDXL quality at roughly 4× the speed. Standard SDXL needs 25–30 steps; this needs 4.
Runtime	PyTorch + diffusers, fp16 on CUDA	Half precision doubles throughput with no visible quality loss at this step count.
Scheduler	Euler, trailing timestep spacing, guidance 0.0	Required by Lightning. Get this wrong and 4-step output is washed out.
Hardware	Cloud GPU — A10G or RTX 4090	Rented per hour, scales horizontally when you need more than 15 batches.
Packaging	Verified zip per batch	One folder per batch, plus a combined archive when you want several at once.

How a day runs

Fifteen batches is 5–7 hours of GPU time, so a full day's 10,500 images completes comfortably within a day. Need more? A second GPU doubles it, a third triples it. Nothing in the pipeline is serial across batches — that's the advantage over running on one computer.

- 01 You send the prompt file (TXT or CSV)
- 02 Queue the batch on the GPU (no manual babysitting)
- 03 Render 700 images in order (~20–30 min)
- 04 Verify every index, then zip (automatic)
- 05 Delivered via Drive or Dropbox link (~3 GB/day)

Throughput, measured not guessed

The engine is verified working end to end on a 4 GB laptop GPU at ~39 s/image — deliberately the slowest hardware I could test on, to prove it degrades gracefully rather than falling over. Production runs on rented GPUs at the rates above. I'll confirm the exact figure on the production GPU during the trial batch rather than ask you to take the table on faith.

GPU	Per image	700 images	10,500 images
RTX 4090	0.6–1.0 s	7–12 min	2–3 h
A10G	1.5–2.5 s	18–30 min	4–7 h
2× A10G	—	—	2–4 h

Built for volume, not demos

- Resume after interruption. If a run stops at image 512, restarting continues from 512. Already-finished images are opened and decoded before being skipped, so a truncated file from a crash is re-rendered rather than delivered.
- Atomic writes. Images are written to a temporary name and renamed only when complete. A crash mid-write cannot leave a corrupt file that looks finished.
- Reproducible. Every image has a recorded seed. If you want number 437 regenerated, I can reproduce or re-roll exactly that one without touching the other 699.
- Per-batch manifest. Each batch ships with a record of prompt, index, seed and render time — useful when you want to trace a specific image back to its line.
- Missing-index detection. Zipping reports any gap. You never receive a folder that's quietly short.

Proposed start

Send me 20–30 prompts from your real file today. I'll return them numbered and zipped so you can check quality and ordering against what you produce now. No charge, no commitment — it's faster than any amount of discussion about whether the system exists.

If the output looks right, we start on full batches immediately. I'd suggest the first paid day be a single batch of 700 so you can confirm the workflow before scaling to 15.

Available to start immediately

Five things worth confirming

01. Are the 15 daily batches 15 different prompt files, or the same 700 prompts rendered 15 times with different seeds? The engine does both — I just need to know which.
02. What resolution and aspect ratio? Default is 1024×1024. Portrait and landscape are both fine, but they change the per-image time.
03. JPEG or PNG? JPEG keeps a day's delivery near 3 GB; PNG would be several times larger to transfer.
04. How would you like delivery — Google Drive, Dropbox, or somewhere else? At this volume, transfer method matters as much as render speed.
05. Any style or negative-prompt defaults to apply across every image, or do the prompts stand alone?

Stated plainly

WHAT IS PROVEN, AND WHAT IS NOT

The engine is built and working — rendering, ordering, verification, resume and packaging are all tested, and it runs today. What I have not yet done is a 10,000-image day on production hardware; the throughput figures above come

from published per-image benchmarks and my own measured timings on slower hardware.

The trial batch settles it either way, at my cost. I'd rather show you a number than claim one.

You asked for someone who already has the system rather than someone who'll build one. Mine is built, it's in version control, and I can run your prompts through it today. Send me a few lines from your file and let's see the output.